

Complex PyTorch for Music Genre Classification

```
In [17]: # Complex pytorch
import torch
import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import DataLoader
from torchvision import datasets, transforms
from complexPyTorch.complexLayers import *
from complexPyTorch.complexFunctions import *

# Plot
import matplotlib.pyplot as plt
import seaborn as sns
import time

# Load Data
import numpy as np
import json
import os
import math
import librosa
import pathlib
from scipy.spatial.distance import cdist
from torch.utils.data import Dataset
from sklearn.model_selection import train_test_split
import random

# MFCCS
from scipy.io import wavfile
import scipy.fftpack as fft
from scipy.signal import get_window
```

```
In [18]: def train(model, device, train_loader, test_loader, optimizer, epoch, met
    model.train()
    total_loss = 0
    correct = 0
    total_samples = len(train_loader.dataset)
    start_time = time.time()

    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        if complexify: data = data.type(torch.complex64)
        if data_fn != None: data = data_fn(data)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)
        loss.backward()
        optimizer.step()
        total_loss += loss.item()
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

    if batch_idx % 10 == 0:
        batch_accuracy = 100. * correct / ((batch_idx + 1) * len(data
        print('Train Epoch: {:3} [{:6}/{:6} ({:3.0f}%)]\tLoss: {:.6f}
```

```

        epoch,
        batch_idx * len(data),
        total_samples,
        100. * batch_idx / len(train_loader),
        loss.item(),
        batch_accuracy)
    )

end_time = time.time()
epoch_times = metrics_dict['epoch_times']
epoch_times.append(end_time - start_time)
epoch_loss = total_loss / len(train_loader)
epoch_accuracy = 100. * correct / total_samples
train_losses = metrics_dict['train_losses']
train_accuracies = metrics_dict['train_accuracies']
train_losses.append(epoch_loss)
train_accuracies.append(epoch_accuracy)
print('Epoch {} - Time: {:.2f}s - Train Loss: {:.6f} - Train Accuracy

# Evaluate on test data
model.eval()
test_loss = 0
correct = 0
with torch.no_grad():
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        if complexify:
            data = data.type(torch.complex64)
        output = model(data)
        test_loss += F.nll_loss(output, target, reduction='sum').item()
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(test_loader.dataset)
test_accuracy = 100. * correct / len(test_loader.dataset)
test_losses = metrics_dict['test_losses']
test_accuracies = metrics_dict['test_accuracies']
test_losses.append(test_loss)
test_accuracies.append(test_accuracy)
print('Test Loss: {:.6f} - Test Accuracy: {:.2f}%\n'.format(test_loss

```

Data Preparation

```

In [19]: DATASET_PATH = "Data/binary_data/train"
SAMPLE_RATE = 22050
TRACK_DURATION = 30 # measured in seconds
SAMPLES_PER_TRACK = SAMPLE_RATE * TRACK_DURATION
BATCH_SIZE = 32
NUM_EPOCHS = 10

```

```

In [20]: genre_list = os.listdir(DATASET_PATH)
if '.DS_Store' in genre_list: genre_list.remove('.DS_Store')
genre_mappings = dict(zip(genre_list, range(len(genre_list))))
print(genre_mappings)

```

```
{'classical': 0, 'rock': 1}
```

MFCCS

```
In [21]: class MusicFeatureExtractor:
    def __init__(self, FFT_size=2048, HOP_SIZE=512, mel_filter_num=13, dct_filter_num=13, epsilon=1e-10):
        self.FFT_size = FFT_size
        self.HOP_SIZE = HOP_SIZE
        self.mel_filter_num = mel_filter_num
        self.dct_filter_num = dct_filter_num
        self.epsilon = 1e-10 # Added to log to avoid log10(0)

    def normalize_audio(self, audio):
        audio = audio / np.max(np.abs(audio))
        return audio

    def frame_audio(self, audio):
        frame_num = int((len(audio) - self.FFT_size) / self.HOP_SIZE) + 1
        frames = np.zeros((frame_num, self.FFT_size))
        for n in range(frame_num):
            frames[n] = audio[n * self.HOP_SIZE: n * self.HOP_SIZE + self.FFT_size]
        return frames

    def freq_to_mel(self, freq):
        return 2595.0 * np.log10(1.0 + freq / 700.0)

    def mel_to_freq(self, mels):
        return 700.0 * (10.0 ** (mels / 2595.0) - 1.0)

    def get_filter_points(self, fmin, fmax, sample_rate):
        fmin_mel = self.freq_to_mel(fmin)
        fmax_mel = self.freq_to_mel(fmax)
        mels = np.linspace(fmin_mel, fmax_mel, num=self.mel_filter_num + 1)
        freqs = self.mel_to_freq(mels)
        return np.floor((self.FFT_size + 1) / sample_rate * freqs).astype(int)

    def get_filters(self, filter_points):
        filters = np.zeros((len(filter_points) - 2, int(self.FFT_size / 2)))
        for n in range(len(filter_points) - 2):
            filters[n, filter_points[n]: filter_points[n + 1]] = np.linspace(0, 1, filter_points[n + 1] - filter_points[n])
            filters[n, filter_points[n + 1]: filter_points[n + 2]] = np.linspace(1, 0, filter_points[n + 2] - filter_points[n + 1])
        return filters

    def dct(self):
        basis = np.empty((self.dct_filter_num, self.mel_filter_num))
        basis[0, :] = 1.0 / np.sqrt(self.mel_filter_num)
        samples = np.arange(1, 2 * self.mel_filter_num, 2) * np.pi / (2.0 * self.FFT_size)
        for i in range(1, self.dct_filter_num):
            basis[i, :] = np.cos(i * samples) * np.sqrt(2.0 / self.mel_filter_num)
        return basis

    def get_mfcc_features(self, audio, sample_rate):
        audio = self.normalize_audio(audio)
        audio_framed = self.frame_audio(audio)
        window = get_window("hann", self.FFT_size, fftbins=True)
        audio_win = audio_framed * window
        audio_winT = np.transpose(audio_win)
        audio_fft = np.empty((int(1 + self.FFT_size // 2), audio_winT.shape[1]))
        for n in range(audio_fft.shape[1]):
            audio_fft[:, n] = fft.fft(audio_winT[:, n], axis=0)[:audio_fft.shape[0]]
```

```

        audio_fft = np.transpose(audio_fft)
        audio_fft = np.square(np.abs(audio_fft))
        freq_min = 0
        freq_high = sample_rate / 2
        filter_points, mel_freqs = self.get_filter_points(freq_min, freq_high)
        filters = self.get_filters(filter_points)
        audio_filtered = np.dot(filters, np.transpose(audio_fft))
        audio_filtered = np.maximum(audio_filtered, self.epsilon) # Replace 0 with epsilon
        audio_log = 10.0 * np.log10(audio_filtered)
        dct_filters = self.dct()
        cepstral_coefficients = np.dot(dct_filters, audio_log)
        return np.array([cepstral_coefficients])

class MusicFeatureExtractorComplex:
    def __init__(self, FFT_size=2048, HOP_SIZE=512, mel_filter_num=13, dct_filter_num=13, epsilon=1e-10):
        self.FFT_size = FFT_size
        self.HOP_SIZE = HOP_SIZE
        self.mel_filter_num = mel_filter_num
        self.dct_filter_num = dct_filter_num
        self.epsilon = 1e-10 # Added to log to avoid log10(0)

    def normalize_audio(self, audio):
        audio = audio / np.max(np.abs(audio))
        return audio

    def frame_audio(self, audio):
        frame_num = int((len(audio) - self.FFT_size) / self.HOP_SIZE) + 1
        frames = np.zeros((frame_num, self.FFT_size))
        for n in range(frame_num):
            frames[n] = audio[n * self.HOP_SIZE: n * self.HOP_SIZE + self.FFT_size]
        return frames

    def freq_to_mel(self, freq):
        return 2595.0 * np.log10(1.0 + freq / 700.0)

    def mel_to_freq(self, mels):
        return 700.0 * (10.0 ** (mels / 2595.0) - 1.0)

    def get_filter_points(self, fmin, fmax, sample_rate):
        fmin_mel = self.freq_to_mel(fmin)
        fmax_mel = self.freq_to_mel(fmax)
        mels = np.linspace(fmin_mel, fmax_mel, num=self.mel_filter_num + 1)
        freqs = self.mel_to_freq(mels)
        return np.floor((self.FFT_size + 1) / sample_rate * freqs).astype(int)

    def get_filters(self, filter_points):
        filters = np.zeros((len(filter_points) - 2, int(self.FFT_size / 2)))
        for n in range(len(filter_points) - 2):
            filters[n, filter_points[n]: filter_points[n + 1]] = np.linspace(0, 1, filter_points[n + 1] - filter_points[n])
            filters[n, filter_points[n + 1]: filter_points[n + 2]] = np.linspace(1, 0, filter_points[n + 2] - filter_points[n + 1])
        return filters

    def dct(self):
        basis = np.empty((self.dct_filter_num, self.mel_filter_num))
        basis[0, :] = 1.0 / np.sqrt(self.mel_filter_num)
        samples = np.arange(1, 2 * self.mel_filter_num, 2) * np.pi / (2.0 * self.mel_filter_num)
        for i in range(1, self.dct_filter_num):
            basis[i, :] = np.cos(i * samples) * np.sqrt(2.0 / self.mel_filter_num)
        return basis

```

```

def get_mfcc_features(self, audio, sample_rate):
    audio = self.normalize_audio(audio)
    audio_framed = self.frame_audio(audio)
    window = get_window("hann", self.FFT_size, fftbins=True)
    audio_win = audio_framed * window
    audio_winT = np.transpose(audio_win)
    audio_fft = np.empty((int(1 + self.FFT_size // 2), audio_winT.shape[1]))
    for n in range(audio_fft.shape[1]):
        audio_fft[:, n] = fft.fft(audio_winT[:, n], axis=0)[:audio_fft.shape[0]]
    audio_fft = np.transpose(audio_fft)
    freq_min = 0
    freq_high = sample_rate / 2
    filter_points, mel_freqs = self.get_filter_points(freq_min, freq_high)
    filters = self.get_filters(filter_points)
    audio_filtered = np.dot(filters, np.transpose(audio_fft))
    audio_filtered[audio_filtered == 0] = self.epsilon # Replace zero
    audio_log = 10.0 * np.log10(audio_filtered)
    dct_filters = self.dct()
    cepstral_coefficients = np.dot(dct_filters, audio_log)
    return np.array([cepstral_coefficients])

```

```

In [22]: class GenreDatasetMFCC(Dataset):

    def __init__(self, train_path, n_fft=2048, hop_length=512, num_segments=10):
        cur_path = pathlib.Path(train_path)
        self.files = []
        for i in list(cur_path.rglob("*.wav")):
            for j in range(num_segments):
                self.files.append([j, i])
        self.samples_per_segment = int(SAMPLES_PER_TRACK / num_segments)
        self.n_fft = n_fft
        self.hop_length = hop_length
        self.num_segments = num_segments
        self.mfcc_extractor = MusicFeatureExtractor(
            FFT_size=n_fft, HOP_SIZE=hop_length, mel_filter_num = mel_filters,
            dct_filter_num = dct_filters)
        self.training = training

    def apply_augmentations(self, signal):
        # Apply augmentations to the audio signal
        if random.random() < 0.5:
            signal = librosa.effects.pitch_shift(signal, sr=SAMPLE_RATE,
            if random.random() < 0.5:
                signal = librosa.effects.time_stretch(signal, rate=random.uniform(0.8, 1.2))
        return signal

    def adjust_shape(self, sequence, max_sequence_length = 126):
        current_length = sequence.shape[2]
        if current_length < max_sequence_length:
            padding = np.zeros((1, 13, max_sequence_length - current_length))
            padded_sequence = np.concatenate((sequence, padding), axis=2)
        else:
            padded_sequence = sequence[:, :, :max_sequence_length]
        return padded_sequence

    def __len__(self):
        return len(self.files)

    def __getitem__(self, idx):
        cur_file = self.files[idx]

```

```

d = cur_file[0]
file_path = cur_file[1]
target = genre_mappings[str(file_path).split("/") [3]]
signal, sample_rate = librosa.load(file_path, sr=SAMPLE_RATE)
start = self.samples_per_segment * d
finish = start + self.samples_per_segment
cur_signal = signal[start:finish]
if self.training: cur_signal = self.apply_augmentations(cur_signal)
cur_mfcc = self.mfcc_extractor.get_mfcc_features(cur_signal, sample_rate)
cur_mfcc = self.adjust_shape(cur_mfcc)
return torch.tensor(cur_mfcc, dtype=torch.float32), target

```

```
class GenreDatasetPhaseMFCC(GenreDatasetMFCC):
```

```

def __init__(self, train_path, n_fft=2048, hop_length=512, num_segments=10, mel_filters=128):
    super().__init__(train_path, n_fft, hop_length, num_segments, mel_filters)
    self.mfcc_extractor = MusicFeatureExtractorComplex(
        FFT_size=n_fft, HOP_SIZE=hop_length, mel_filter_num = mel_filters)

def __getitem__(self, idx):
    cur_file = self.files[idx]
    d = cur_file[0]
    file_path = cur_file[1]
    target = genre_mappings[str(file_path).split("/") [3]]
    signal, sample_rate = librosa.load(file_path, sr=SAMPLE_RATE)
    start = self.samples_per_segment * d
    finish = start + self.samples_per_segment
    cur_signal = signal[start:finish]
    if self.training: cur_signal = self.apply_augmentations(cur_signal)
    cur_mfcc = self.mfcc_extractor.get_mfcc_features(cur_signal, sample_rate)
    cur_mfcc = self.adjust_shape(cur_mfcc)
    return torch.tensor(cur_mfcc, dtype=torch.complex64), target

```

1. No phase data

```

In [30]: train_dataset = GenreDatasetMFCC("Data/binary_data/train/", n_fft=2048, hop_length=512, num_segments=10, mel_filters=128)
test_dataset = GenreDatasetMFCC("Data/binary_data/test/", n_fft=2048, hop_length=512, num_segments=10, mel_filters=128)
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_dataset, shuffle=False)

```

```
In [31]: class RealNet(nn.Module):
```

```

def __init__(self):
    super(RealNet, self).__init__()
    self.conv1 = nn.Conv2d(1, 10, 2, 1)
    self.bn = nn.BatchNorm2d(10)
    self.conv2 = nn.Conv2d(10, 20, 2, 1)
    self.fc1 = nn.Linear(30*2*20, 500)
    self.fc2 = nn.Linear(500, 2)

def forward(self, x):
    x = self.conv1(x)
    x = F.relu(x)
    x = F.max_pool2d(x, 2, 2)
    x = self.bn(x)
    x = self.conv2(x)
    x = F.relu(x)
    x = F.max_pool2d(x, 2, 2)

```

```
x = x.view(-1,30*2*20)
x = self.fc1(x)
x = F.relu(x)
x = self.fc2(x)
x = x.abs()
x = F.log_softmax(x, dim=1)
return x

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = RealNet().to(device)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

metrics_dict_e1 = {
    'epoch_times': [],
    'train_losses': [],
    'train_accuracies': [],
    'test_losses': [],
    'test_accuracies': []
}

for epoch in range(NUM_EPOCHS):
    train(model,
          device,
          train_loader,
          test_loader,
          optimizer,
          epoch,
          metrics_dict_e1,
          complexify = False)

print("-"*100)
print("-"*100)
print("FINAL RESULTS:")
print("-"*100)
for key, value in metrics_dict_e1.items():
    print(f'{key}: {value}')
```

Train Epoch: 0 [0/ 1680 (0%)] Loss: 0.715544 Accuracy: 34.38%
Train Epoch: 0 [320/ 1680 (19%)] Loss: 0.425783 Accuracy: 57.95%
Train Epoch: 0 [640/ 1680 (38%)] Loss: 0.214916 Accuracy: 72.17%
Train Epoch: 0 [960/ 1680 (58%)] Loss: 0.330717 Accuracy: 77.92%
Train Epoch: 0 [1280/ 1680 (77%)] Loss: 0.253875 Accuracy: 80.72%
Train Epoch: 0 [1600/ 1680 (96%)] Loss: 0.151264 Accuracy: 82.60%
Epoch 0 – Time: 37.52s – Train Loss: 0.361252 – Train Accuracy: 81.90%
Test Loss: 0.184620 – Test Accuracy: 95.00%

Train Epoch: 1 [0/ 1680 (0%)] Loss: 0.116859 Accuracy: 93.75%
Train Epoch: 1 [320/ 1680 (19%)] Loss: 0.181167 Accuracy: 91.48%
Train Epoch: 1 [640/ 1680 (38%)] Loss: 0.205668 Accuracy: 91.52%
Train Epoch: 1 [960/ 1680 (58%)] Loss: 0.297967 Accuracy: 91.94%
Train Epoch: 1 [1280/ 1680 (77%)] Loss: 0.142673 Accuracy: 91.92%
Train Epoch: 1 [1600/ 1680 (96%)] Loss: 0.586391 Accuracy: 91.42%
Epoch 1 – Time: 51.89s – Train Loss: 0.233470 – Train Accuracy: 90.65%
Test Loss: 0.176717 – Test Accuracy: 96.25%

Train Epoch: 2 [0/ 1680 (0%)] Loss: 0.191501 Accuracy: 93.75%
Train Epoch: 2 [320/ 1680 (19%)] Loss: 0.473941 Accuracy: 92.61%
Train Epoch: 2 [640/ 1680 (38%)] Loss: 0.150896 Accuracy: 91.52%
Train Epoch: 2 [960/ 1680 (58%)] Loss: 0.240908 Accuracy: 91.43%
Train Epoch: 2 [1280/ 1680 (77%)] Loss: 0.205249 Accuracy: 91.16%
Train Epoch: 2 [1600/ 1680 (96%)] Loss: 0.145471 Accuracy: 91.24%
Epoch 2 – Time: 39.21s – Train Loss: 0.231075 – Train Accuracy: 90.42%
Test Loss: 0.192560 – Test Accuracy: 96.25%

Train Epoch: 3 [0/ 1680 (0%)] Loss: 0.165502 Accuracy: 96.88%
Train Epoch: 3 [320/ 1680 (19%)] Loss: 0.410278 Accuracy: 91.48%
Train Epoch: 3 [640/ 1680 (38%)] Loss: 0.271823 Accuracy: 92.41%
Train Epoch: 3 [960/ 1680 (58%)] Loss: 0.213464 Accuracy: 93.45%
Train Epoch: 3 [1280/ 1680 (77%)] Loss: 0.102448 Accuracy: 92.76%
Train Epoch: 3 [1600/ 1680 (96%)] Loss: 0.182957 Accuracy: 91.79%
Epoch 3 – Time: 41.20s – Train Loss: 0.215985 – Train Accuracy: 90.65%
Test Loss: 0.199063 – Test Accuracy: 94.38%

Train Epoch: 4 [0/ 1680 (0%)] Loss: 0.283567 Accuracy: 90.62%
Train Epoch: 4 [320/ 1680 (19%)] Loss: 0.214640 Accuracy: 92.05%
Train Epoch: 4 [640/ 1680 (38%)] Loss: 0.326049 Accuracy: 93.01%
Train Epoch: 4 [960/ 1680 (58%)] Loss: 0.086493 Accuracy: 93.55%
Train Epoch: 4 [1280/ 1680 (77%)] Loss: 0.070953 Accuracy: 93.52%
Train Epoch: 4 [1600/ 1680 (96%)] Loss: 0.176417 Accuracy: 93.44%
Epoch 4 – Time: 38.53s – Train Loss: 0.190518 – Train Accuracy: 92.68%
Test Loss: 0.184343 – Test Accuracy: 95.31%

Train Epoch: 5 [0/ 1680 (0%)] Loss: 0.078587 Accuracy: 96.88%
Train Epoch: 5 [320/ 1680 (19%)] Loss: 0.150306 Accuracy: 92.33%
Train Epoch: 5 [640/ 1680 (38%)] Loss: 0.165681 Accuracy: 91.82%
Train Epoch: 5 [960/ 1680 (58%)] Loss: 0.112581 Accuracy: 92.84%
Train Epoch: 5 [1280/ 1680 (77%)] Loss: 0.181829 Accuracy: 92.68%
Train Epoch: 5 [1600/ 1680 (96%)] Loss: 0.231670 Accuracy: 93.01%
Epoch 5 – Time: 40.21s – Train Loss: 0.179768 – Train Accuracy: 92.08%
Test Loss: 0.237013 – Test Accuracy: 95.00%

Train Epoch: 6 [0/ 1680 (0%)] Loss: 0.332018 Accuracy: 81.25%
Train Epoch: 6 [320/ 1680 (19%)] Loss: 0.118397 Accuracy: 91.19%
Train Epoch: 6 [640/ 1680 (38%)] Loss: 0.219945 Accuracy: 92.56%
Train Epoch: 6 [960/ 1680 (58%)] Loss: 0.310991 Accuracy: 92.84%
Train Epoch: 6 [1280/ 1680 (77%)] Loss: 0.037779 Accuracy: 93.67%
Train Epoch: 6 [1600/ 1680 (96%)] Loss: 0.269485 Accuracy: 93.63%

Epoch 6 – Time: 52.36s – Train Loss: 0.172697 – Train Accuracy: 92.74%
 Test Loss: 0.295163 – Test Accuracy: 91.88%

Train Epoch: 7 [0/ 1680 (0%)] Loss: 0.177706 Accuracy: 90.62%
 Train Epoch: 7 [320/ 1680 (19%)] Loss: 0.142058 Accuracy: 91.76%
 Train Epoch: 7 [640/ 1680 (38%)] Loss: 0.085112 Accuracy: 93.60%
 Train Epoch: 7 [960/ 1680 (58%)] Loss: 0.366243 Accuracy: 94.25%
 Train Epoch: 7 [1280/ 1680 (77%)] Loss: 0.244307 Accuracy: 93.90%
 Train Epoch: 7 [1600/ 1680 (96%)] Loss: 0.198373 Accuracy: 93.38%
 Epoch 7 – Time: 36.99s – Train Loss: 0.185385 – Train Accuracy: 92.44%
 Test Loss: 0.174413 – Test Accuracy: 96.88%

Train Epoch: 8 [0/ 1680 (0%)] Loss: 0.159421 Accuracy: 96.88%
 Train Epoch: 8 [320/ 1680 (19%)] Loss: 0.228357 Accuracy: 91.76%
 Train Epoch: 8 [640/ 1680 (38%)] Loss: 0.124323 Accuracy: 93.30%
 Train Epoch: 8 [960/ 1680 (58%)] Loss: 0.052267 Accuracy: 93.25%
 Train Epoch: 8 [1280/ 1680 (77%)] Loss: 0.523930 Accuracy: 92.99%
 Train Epoch: 8 [1600/ 1680 (96%)] Loss: 0.190568 Accuracy: 93.44%
 Epoch 8 – Time: 35.39s – Train Loss: 0.179532 – Train Accuracy: 92.50%
 Test Loss: 0.201300 – Test Accuracy: 92.50%

Train Epoch: 9 [0/ 1680 (0%)] Loss: 0.179819 Accuracy: 93.75%
 Train Epoch: 9 [320/ 1680 (19%)] Loss: 0.231372 Accuracy: 94.03%
 Train Epoch: 9 [640/ 1680 (38%)] Loss: 0.117419 Accuracy: 95.09%
 Train Epoch: 9 [960/ 1680 (58%)] Loss: 0.182112 Accuracy: 95.26%
 Train Epoch: 9 [1280/ 1680 (77%)] Loss: 0.084204 Accuracy: 94.97%
 Train Epoch: 9 [1600/ 1680 (96%)] Loss: 0.281102 Accuracy: 94.49%
 Epoch 9 – Time: 36.27s – Train Loss: 0.138495 – Train Accuracy: 93.69%
 Test Loss: 0.196841 – Test Accuracy: 96.56%

 FINAL RESULTS:


```
epoch_times: [37.51734495162964, 51.892842054367065, 39.20804286003113, 4
1.200674057006836, 38.53182005882263, 40.20562291145325, 52.3609809875488
3, 36.98588514328003, 35.39342212677002, 36.267234086990356]
train_losses: [0.3612523521654881, 0.23346950615254733, 0.2310745389415667
6, 0.2159847513271066, 0.19051814960459104, 0.17976838558052594, 0.1726974
0708602163, 0.18538537392249474, 0.17953154215445885, 0.1384946137953263]
train_accuracies: [81.9047619047619, 90.6547619047619, 90.41666666666667,
90.6547619047619, 92.67857142857143, 92.08333333333333, 92.73809523809524,
92.44047619047619, 92.5, 93.69047619047619]
test_losses: [0.18461954668164254, 0.1767165631055832, 0.1925600398331880
6, 0.19906260073184967, 0.18434346728026868, 0.23701292220503092, 0.295162
65785787255, 0.17441322961822153, 0.20130045646801592, 0.1968405670020729
2]
test_accuracies: [95.0, 96.25, 96.25, 94.375, 95.3125, 95.0, 91.875, 96.87
5, 92.5, 96.5625]
```

In [32]: `class ComplexNet(nn.Module):`

```
def __init__(self):
    super(ComplexNet, self).__init__()
    self.conv1 = ComplexConv2d(1, 10, 2, 1)
    self.bn = ComplexBatchNorm2d(10)
    self.conv2 = ComplexConv2d(10, 20, 2, 1)
```

```

self.fc1 = ComplexLinear(30*2*20, 500)
self.fc2 = ComplexLinear(500, 2)

def forward(self,x):
    x = self.conv1(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = self.bn(x)
    x = self.conv2(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = x.view(-1,30*2*20)
    x = self.fc1(x)
    x = complex_relu(x)
    x = self.fc2(x)
    x = x.abs()
    x = F.log_softmax(x, dim=1)
    return x

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = ComplexNet().to(device)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

metrics_dict_e2 = {
    'epoch_times': [],
    'train_losses': [],
    'train_accuracies': [],
    'test_losses': [],
    'test_accuracies': []
}

for epoch in range(NUM_EPOCHS):
    train(model,
          device,
          train_loader,
          test_loader,
          optimizer,
          epoch,
          metrics_dict_e2)

print("-"*100)
print("-"*100)
print("FINAL RESULTS:")
print("-"*100)
for key, value in metrics_dict_e2.items():
    print(f'{key}: {value}')

```

Train Epoch: 0 [0/ 1680 (0%)] Loss: 0.679613 Accuracy: 62.50%
Train Epoch: 0 [320/ 1680 (19%)] Loss: 0.446472 Accuracy: 73.58%
Train Epoch: 0 [640/ 1680 (38%)] Loss: 0.108056 Accuracy: 80.80%
Train Epoch: 0 [960/ 1680 (58%)] Loss: 0.184048 Accuracy: 84.48%
Train Epoch: 0 [1280/ 1680 (77%)] Loss: 0.135176 Accuracy: 87.04%
Train Epoch: 0 [1600/ 1680 (96%)] Loss: 0.307315 Accuracy: 87.19%
Epoch 0 – Time: 42.76s – Train Loss: 0.312190 – Train Accuracy: 86.43%
Test Loss: 0.181000 – Test Accuracy: 94.69%

Train Epoch: 1 [0/ 1680 (0%)] Loss: 0.220294 Accuracy: 93.75%
Train Epoch: 1 [320/ 1680 (19%)] Loss: 0.210621 Accuracy: 93.75%
Train Epoch: 1 [640/ 1680 (38%)] Loss: 0.052198 Accuracy: 93.90%
Train Epoch: 1 [960/ 1680 (58%)] Loss: 0.055498 Accuracy: 93.95%
Train Epoch: 1 [1280/ 1680 (77%)] Loss: 0.133853 Accuracy: 93.67%
Train Epoch: 1 [1600/ 1680 (96%)] Loss: 0.059262 Accuracy: 93.93%
Epoch 1 – Time: 43.32s – Train Loss: 0.173853 – Train Accuracy: 93.04%
Test Loss: 0.183872 – Test Accuracy: 95.31%

Train Epoch: 2 [0/ 1680 (0%)] Loss: 0.021537 Accuracy: 100.00%
Train Epoch: 2 [320/ 1680 (19%)] Loss: 0.261391 Accuracy: 94.32%
Train Epoch: 2 [640/ 1680 (38%)] Loss: 0.105262 Accuracy: 93.75%
Train Epoch: 2 [960/ 1680 (58%)] Loss: 0.164997 Accuracy: 93.95%
Train Epoch: 2 [1280/ 1680 (77%)] Loss: 0.156987 Accuracy: 93.67%
Train Epoch: 2 [1600/ 1680 (96%)] Loss: 0.117551 Accuracy: 93.26%
Epoch 2 – Time: 43.32s – Train Loss: 0.181060 – Train Accuracy: 92.38%
Test Loss: 0.294938 – Test Accuracy: 92.19%

Train Epoch: 3 [0/ 1680 (0%)] Loss: 0.141607 Accuracy: 96.88%
Train Epoch: 3 [320/ 1680 (19%)] Loss: 0.124201 Accuracy: 93.47%
Train Epoch: 3 [640/ 1680 (38%)] Loss: 0.124381 Accuracy: 92.71%
Train Epoch: 3 [960/ 1680 (58%)] Loss: 0.291106 Accuracy: 92.44%
Train Epoch: 3 [1280/ 1680 (77%)] Loss: 0.289438 Accuracy: 92.68%
Train Epoch: 3 [1600/ 1680 (96%)] Loss: 0.119186 Accuracy: 92.71%
Epoch 3 – Time: 42.87s – Train Loss: 0.175521 – Train Accuracy: 91.90%
Test Loss: 0.160474 – Test Accuracy: 95.94%

Train Epoch: 4 [0/ 1680 (0%)] Loss: 0.117496 Accuracy: 100.00%
Train Epoch: 4 [320/ 1680 (19%)] Loss: 0.089305 Accuracy: 94.60%
Train Epoch: 4 [640/ 1680 (38%)] Loss: 0.145823 Accuracy: 94.94%
Train Epoch: 4 [960/ 1680 (58%)] Loss: 0.103748 Accuracy: 95.26%
Train Epoch: 4 [1280/ 1680 (77%)] Loss: 0.242247 Accuracy: 94.97%
Train Epoch: 4 [1600/ 1680 (96%)] Loss: 0.122512 Accuracy: 94.85%
Epoch 4 – Time: 44.11s – Train Loss: 0.134570 – Train Accuracy: 93.87%
Test Loss: 0.228681 – Test Accuracy: 90.94%

Train Epoch: 5 [0/ 1680 (0%)] Loss: 0.101651 Accuracy: 96.88%
Train Epoch: 5 [320/ 1680 (19%)] Loss: 0.208671 Accuracy: 94.89%
Train Epoch: 5 [640/ 1680 (38%)] Loss: 0.072090 Accuracy: 94.94%
Train Epoch: 5 [960/ 1680 (58%)] Loss: 0.259198 Accuracy: 94.35%
Train Epoch: 5 [1280/ 1680 (77%)] Loss: 0.106292 Accuracy: 94.13%
Train Epoch: 5 [1600/ 1680 (96%)] Loss: 0.110741 Accuracy: 94.36%
Epoch 5 – Time: 43.58s – Train Loss: 0.141572 – Train Accuracy: 93.45%
Test Loss: 0.118991 – Test Accuracy: 97.19%

Train Epoch: 6 [0/ 1680 (0%)] Loss: 0.238775 Accuracy: 84.38%
Train Epoch: 6 [320/ 1680 (19%)] Loss: 0.122788 Accuracy: 94.03%
Train Epoch: 6 [640/ 1680 (38%)] Loss: 0.012993 Accuracy: 95.98%
Train Epoch: 6 [960/ 1680 (58%)] Loss: 0.094011 Accuracy: 95.87%
Train Epoch: 6 [1280/ 1680 (77%)] Loss: 0.145716 Accuracy: 95.58%
Train Epoch: 6 [1600/ 1680 (96%)] Loss: 0.544419 Accuracy: 95.47%

Epoch 6 – Time: 58.67s – Train Loss: 0.129391 – Train Accuracy: 94.46%
 Test Loss: 0.152756 – Test Accuracy: 96.56%

Train Epoch: 7 [0/ 1680 (0%)] Loss: 0.110688 Accuracy: 96.88%
 Train Epoch: 7 [320/ 1680 (19%)] Loss: 0.115284 Accuracy: 95.45%
 Train Epoch: 7 [640/ 1680 (38%)] Loss: 0.113117 Accuracy: 95.39%
 Train Epoch: 7 [960/ 1680 (58%)] Loss: 0.078003 Accuracy: 95.97%
 Train Epoch: 7 [1280/ 1680 (77%)] Loss: 0.191643 Accuracy: 96.04%
 Train Epoch: 7 [1600/ 1680 (96%)] Loss: 0.159519 Accuracy: 95.96%
 Epoch 7 – Time: 45.36s – Train Loss: 0.102312 – Train Accuracy: 95.00%
 Test Loss: 0.215539 – Test Accuracy: 95.00%

Train Epoch: 8 [0/ 1680 (0%)] Loss: 0.151010 Accuracy: 96.88%
 Train Epoch: 8 [320/ 1680 (19%)] Loss: 0.082022 Accuracy: 94.03%
 Train Epoch: 8 [640/ 1680 (38%)] Loss: 0.019679 Accuracy: 95.54%
 Train Epoch: 8 [960/ 1680 (58%)] Loss: 0.090209 Accuracy: 94.66%
 Train Epoch: 8 [1280/ 1680 (77%)] Loss: 0.244851 Accuracy: 94.59%
 Train Epoch: 8 [1600/ 1680 (96%)] Loss: 0.130030 Accuracy: 94.24%
 Epoch 8 – Time: 49.18s – Train Loss: 0.145277 – Train Accuracy: 93.39%
 Test Loss: 0.354122 – Test Accuracy: 86.56%

Train Epoch: 9 [0/ 1680 (0%)] Loss: 0.343555 Accuracy: 84.38%
 Train Epoch: 9 [320/ 1680 (19%)] Loss: 0.080193 Accuracy: 94.89%
 Train Epoch: 9 [640/ 1680 (38%)] Loss: 0.105633 Accuracy: 95.09%
 Train Epoch: 9 [960/ 1680 (58%)] Loss: 0.054809 Accuracy: 94.76%
 Train Epoch: 9 [1280/ 1680 (77%)] Loss: 0.095516 Accuracy: 95.05%
 Train Epoch: 9 [1600/ 1680 (96%)] Loss: 0.084872 Accuracy: 95.22%
 Epoch 9 – Time: 44.47s – Train Loss: 0.117733 – Train Accuracy: 94.29%
 Test Loss: 0.195516 – Test Accuracy: 95.00%

FINAL RESULTS:

```
epoch_times: [42.76304793357849, 43.31869101524353, 43.31773591041565, 42.86594581604004, 44.11113405227661, 43.584967851638794, 58.66502928733826, 45.360710859298706, 49.17667317390442, 44.471909046173096]
train_losses: [0.31219020901391137, 0.17385263096254605, 0.18106022778038794, 0.17552123677272063, 0.13456989785369772, 0.14157151903670567, 0.12939078459301248, 0.10231199516699864, 0.14527743521074837, 0.11773306253151251]
train_accuracies: [86.42857142857143, 93.03571428571429, 92.38095238095238, 91.9047619047619, 93.86904761904762, 93.45238095238095, 94.46428571428571, 95.0, 93.39285714285714, 94.28571428571429]
test_losses: [0.18100014608353376, 0.18387161334976554, 0.29493764862418176, 0.16047407081350684, 0.22868142360821367, 0.11899070171639323, 0.15275583432521672, 0.2155387081365916, 0.3541224246728234, 0.19551602460269352]
test_accuracies: [94.6875, 95.3125, 92.1875, 95.9375, 90.9375, 97.1875, 96.5625, 95.0, 86.5625, 95.0]
```

```
In [33]: train_dataset = GenreDatasetPhaseMFCC("Data/binary_data/train/", n_fft=2048)
test_dataset = GenreDatasetPhaseMFCC("Data/binary_data/test/", n_fft=2048)
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_dataset, shuffle=False)

class ComplexNet(nn.Module):
```

```

def __init__(self):
    super(ComplexNet, self).__init__()
    self.conv1 = ComplexConv2d(1, 10, 2, 1)
    self.bn = ComplexBatchNorm2d(10)
    self.conv2 = ComplexConv2d(10, 20, 2, 1)
    self.fc1 = ComplexLinear(30*2*20, 500)
    self.fc2 = ComplexLinear(500, 2)

def forward(self, x):
    x = self.conv1(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = self.bn(x)
    x = self.conv2(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = x.view(-1, 30*2*20)
    x = self.fc1(x)
    x = complex_relu(x)
    x = self.fc2(x)
    x = x.abs()
    x = F.log_softmax(x, dim=1)
    return x

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = ComplexNet().to(device)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

metrics_dict_e3 = {
    'epoch_times': [],
    'train_losses': [],
    'train_accuracies': [],
    'test_losses': [],
    'test_accuracies': []
}

for epoch in range(NUM_EPOCHS):
    train(model,
          device,
          train_loader,
          test_loader,
          optimizer,
          epoch,
          metrics_dict_e3)

print("-"*100)
print("-"*100)
print("FINAL RESULTS:")
print("-"*100)
for key, value in metrics_dict_e3.items():
    print(f'{key}: {value}')

```

Train Epoch: 0 [0/ 1680 (0%)] Loss: 0.735650 Accuracy: 53.12%
Train Epoch: 0 [320/ 1680 (19%)] Loss: 0.868494 Accuracy: 51.14%
Train Epoch: 0 [640/ 1680 (38%)] Loss: 0.549072 Accuracy: 55.95%
Train Epoch: 0 [960/ 1680 (58%)] Loss: 0.466943 Accuracy: 62.40%
Train Epoch: 0 [1280/ 1680 (77%)] Loss: 0.244168 Accuracy: 68.29%
Train Epoch: 0 [1600/ 1680 (96%)] Loss: 0.429500 Accuracy: 71.69%
Epoch 0 – Time: 73.58s – Train Loss: 0.603171 – Train Accuracy: 71.31%
Test Loss: 0.336265 – Test Accuracy: 87.50%

Train Epoch: 1 [0/ 1680 (0%)] Loss: 0.425153 Accuracy: 78.12%
Train Epoch: 1 [320/ 1680 (19%)] Loss: 0.595793 Accuracy: 77.27%
Train Epoch: 1 [640/ 1680 (38%)] Loss: 0.317432 Accuracy: 82.44%
Train Epoch: 1 [960/ 1680 (58%)] Loss: 0.321401 Accuracy: 82.66%
Train Epoch: 1 [1280/ 1680 (77%)] Loss: 0.155262 Accuracy: 83.08%
Train Epoch: 1 [1600/ 1680 (96%)] Loss: 0.352507 Accuracy: 82.60%
Epoch 1 – Time: 51.37s – Train Loss: 0.424381 – Train Accuracy: 81.61%
Test Loss: 0.602541 – Test Accuracy: 66.56%

Train Epoch: 2 [0/ 1680 (0%)] Loss: 0.850426 Accuracy: 50.00%
Train Epoch: 2 [320/ 1680 (19%)] Loss: 0.239433 Accuracy: 78.98%
Train Epoch: 2 [640/ 1680 (38%)] Loss: 0.500215 Accuracy: 82.29%
Train Epoch: 2 [960/ 1680 (58%)] Loss: 0.218858 Accuracy: 83.97%
Train Epoch: 2 [1280/ 1680 (77%)] Loss: 0.332864 Accuracy: 85.59%
Train Epoch: 2 [1600/ 1680 (96%)] Loss: 0.262290 Accuracy: 86.95%
Epoch 2 – Time: 57.94s – Train Loss: 0.305707 – Train Accuracy: 86.13%
Test Loss: 0.280999 – Test Accuracy: 90.31%

Train Epoch: 3 [0/ 1680 (0%)] Loss: 0.145187 Accuracy: 93.75%
Train Epoch: 3 [320/ 1680 (19%)] Loss: 0.264619 Accuracy: 88.35%
Train Epoch: 3 [640/ 1680 (38%)] Loss: 0.230185 Accuracy: 87.65%
Train Epoch: 3 [960/ 1680 (58%)] Loss: 0.407401 Accuracy: 87.80%
Train Epoch: 3 [1280/ 1680 (77%)] Loss: 0.211505 Accuracy: 88.49%
Train Epoch: 3 [1600/ 1680 (96%)] Loss: 0.181106 Accuracy: 89.09%
Epoch 3 – Time: 47.14s – Train Loss: 0.285976 – Train Accuracy: 88.33%
Test Loss: 0.276645 – Test Accuracy: 91.56%

Train Epoch: 4 [0/ 1680 (0%)] Loss: 0.497322 Accuracy: 87.50%
Train Epoch: 4 [320/ 1680 (19%)] Loss: 0.536946 Accuracy: 87.50%
Train Epoch: 4 [640/ 1680 (38%)] Loss: 0.299088 Accuracy: 88.54%
Train Epoch: 4 [960/ 1680 (58%)] Loss: 0.295729 Accuracy: 87.90%
Train Epoch: 4 [1280/ 1680 (77%)] Loss: 0.254625 Accuracy: 88.34%
Train Epoch: 4 [1600/ 1680 (96%)] Loss: 0.490942 Accuracy: 88.30%
Epoch 4 – Time: 51.99s – Train Loss: 0.285149 – Train Accuracy: 87.50%
Test Loss: 0.246197 – Test Accuracy: 90.94%

Train Epoch: 5 [0/ 1680 (0%)] Loss: 0.327977 Accuracy: 84.38%
Train Epoch: 5 [320/ 1680 (19%)] Loss: 0.147034 Accuracy: 91.76%
Train Epoch: 5 [640/ 1680 (38%)] Loss: 0.187032 Accuracy: 90.48%
Train Epoch: 5 [960/ 1680 (58%)] Loss: 0.338331 Accuracy: 90.22%
Train Epoch: 5 [1280/ 1680 (77%)] Loss: 0.390736 Accuracy: 88.26%
Train Epoch: 5 [1600/ 1680 (96%)] Loss: 0.275194 Accuracy: 88.05%
Epoch 5 – Time: 51.88s – Train Loss: 0.301933 – Train Accuracy: 87.32%
Test Loss: 0.271775 – Test Accuracy: 89.06%

Train Epoch: 6 [0/ 1680 (0%)] Loss: 0.331382 Accuracy: 84.38%
Train Epoch: 6 [320/ 1680 (19%)] Loss: 0.304538 Accuracy: 90.62%
Train Epoch: 6 [640/ 1680 (38%)] Loss: 0.093887 Accuracy: 91.37%
Train Epoch: 6 [960/ 1680 (58%)] Loss: 0.245574 Accuracy: 91.03%
Train Epoch: 6 [1280/ 1680 (77%)] Loss: 0.315051 Accuracy: 89.10%
Train Epoch: 6 [1600/ 1680 (96%)] Loss: 0.260195 Accuracy: 89.28%

Epoch 6 – Time: 48.36s – Train Loss: 0.282665 – Train Accuracy: 88.33%
 Test Loss: 0.210059 – Test Accuracy: 90.94%

Train Epoch: 7 [0/ 1680 (0%)] Loss: 0.095637 Accuracy: 96.88%
 Train Epoch: 7 [320/ 1680 (19%)] Loss: 0.291555 Accuracy: 90.34%
 Train Epoch: 7 [640/ 1680 (38%)] Loss: 0.313008 Accuracy: 89.73%
 Train Epoch: 7 [960/ 1680 (58%)] Loss: 0.348511 Accuracy: 89.82%
 Train Epoch: 7 [1280/ 1680 (77%)] Loss: 0.470430 Accuracy: 89.33%
 Train Epoch: 7 [1600/ 1680 (96%)] Loss: 0.302118 Accuracy: 89.15%
 Epoch 7 – Time: 48.26s – Train Loss: 0.269804 – Train Accuracy: 88.45%
 Test Loss: 0.201279 – Test Accuracy: 92.81%

Train Epoch: 8 [0/ 1680 (0%)] Loss: 0.291109 Accuracy: 87.50%
 Train Epoch: 8 [320/ 1680 (19%)] Loss: 0.220785 Accuracy: 89.49%
 Train Epoch: 8 [640/ 1680 (38%)] Loss: 0.301213 Accuracy: 88.54%
 Train Epoch: 8 [960/ 1680 (58%)] Loss: 0.490905 Accuracy: 88.81%
 Train Epoch: 8 [1280/ 1680 (77%)] Loss: 0.228495 Accuracy: 89.25%
 Train Epoch: 8 [1600/ 1680 (96%)] Loss: 0.160624 Accuracy: 89.71%
 Epoch 8 – Time: 48.60s – Train Loss: 0.256629 – Train Accuracy: 88.99%
 Test Loss: 0.293889 – Test Accuracy: 90.00%

Train Epoch: 9 [0/ 1680 (0%)] Loss: 0.375799 Accuracy: 81.25%
 Train Epoch: 9 [320/ 1680 (19%)] Loss: 0.247870 Accuracy: 89.49%
 Train Epoch: 9 [640/ 1680 (38%)] Loss: 0.087902 Accuracy: 90.62%
 Train Epoch: 9 [960/ 1680 (58%)] Loss: 0.392372 Accuracy: 90.02%
 Train Epoch: 9 [1280/ 1680 (77%)] Loss: 0.262468 Accuracy: 88.95%
 Train Epoch: 9 [1600/ 1680 (96%)] Loss: 0.162650 Accuracy: 88.91%
 Epoch 9 – Time: 48.17s – Train Loss: 0.255111 – Train Accuracy: 88.21%
 Test Loss: 0.277363 – Test Accuracy: 91.56%

 FINAL RESULTS:

epoch_times: [73.57843279838562, 51.3704469203949, 57.93836307525635, 47.135130882263184, 51.98855710029602, 51.87985110282898, 48.36233329772949, 48.26110076904297, 48.60330104827881, 48.16680717468262]
 train_losses: [0.603170841645736, 0.4243808463215828, 0.30570677667856216, 0.285976423523747, 0.2851490078923794, 0.30193280887145263, 0.2826654441081561, 0.2698039590882567, 0.25662878531819355, 0.25511147984518456]
 train_accuracies: [71.30952380952381, 81.60714285714286, 86.13095238095238, 88.33333333333333, 87.5, 87.32142857142857, 88.33333333333333, 88.45238095238095, 88.98809523809524, 88.21428571428571]
 test_losses: [0.3362645523622632, 0.6025413427501917, 0.2809987593907863, 0.2766452480107546, 0.24619730543345214, 0.27177466712892057, 0.21005938109010458, 0.20127928256988525, 0.29388898848555983, 0.2773632241412997]
 test_accuracies: [87.5, 66.5625, 90.3125, 91.5625, 90.9375, 89.0625, 90.9375, 92.8125, 90.0, 91.5625]

Plots

```
In [35]: # Data for the four scenarios
data = {
    "Magnitude Only (Real Net)": metrics_dict_e1,
    "Magnitude Only (Complex Net)": metrics_dict_e2,
```

```

    "Magnitude and Phase (Complex Net)": metrics_dict_e3
}

# Data for plotting
epochs = range(1, 11)
colors = ['b', 'g', 'r', 'm', 'y']
scenarios = list(data.keys())

fig, axes = plt.subplots(2, 1, figsize=(10, 10))

for i, scenario in enumerate(scenarios):
    axes[0].plot(epochs, data[scenario]["train accuracies"], label=scenario)

axes[0].set_title("Train Accuracy")
axes[0].set_xlabel("Epochs")
axes[0].set_ylabel("Train Accuracy")
axes[0].legend()

for i, scenario in enumerate(scenarios):
    axes[1].plot(epochs, data[scenario]["test accuracies"], label=scenario)

axes[1].set_title("Test Accuracy")
axes[1].set_xlabel("Epochs")
axes[1].set_ylabel("Test Accuracy")
axes[1].legend()

plt.tight_layout()
plt.show()

fig, axes = plt.subplots(2, 1, figsize=(10, 10))

for i, scenario in enumerate(scenarios):
    axes[0].plot(epochs, data[scenario]["train losses"], label=scenario, color=colors[i])

axes[0].set_title("Train Loss")
axes[0].set_xlabel("Epochs")
axes[0].set_ylabel("Train Loss")
axes[0].legend()


for i, scenario in enumerate(scenarios):
    axes[1].plot(epochs, data[scenario]["test losses"], label=scenario, color=colors[i])

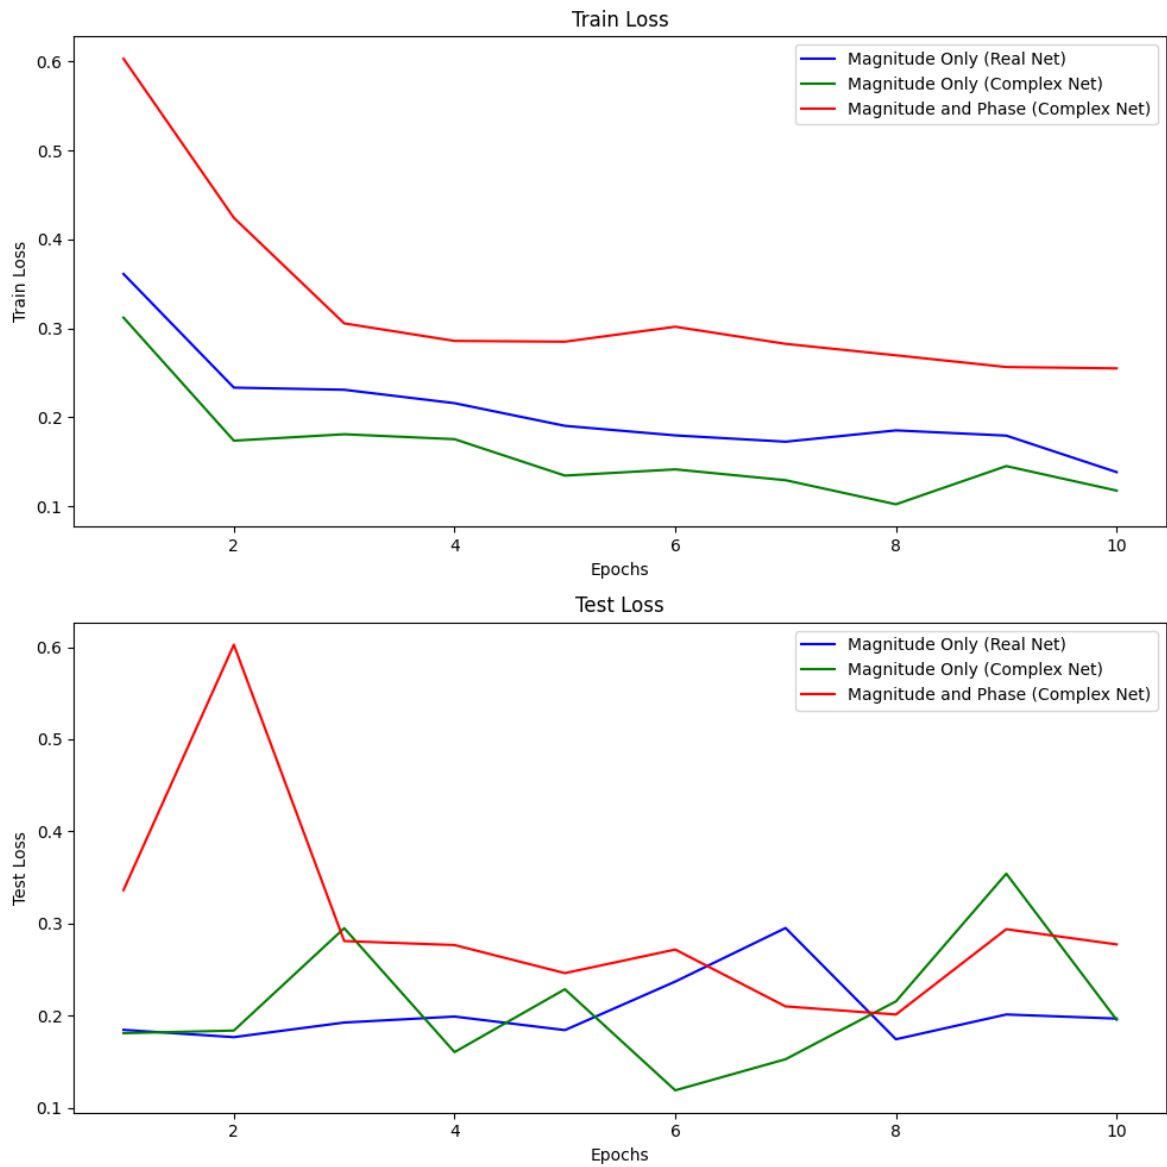
axes[1].set_title("Test Loss")
axes[1].set_xlabel("Epochs")
axes[1].set_ylabel("Test Loss")
axes[1].legend()

plt.tight_layout()
plt.show()

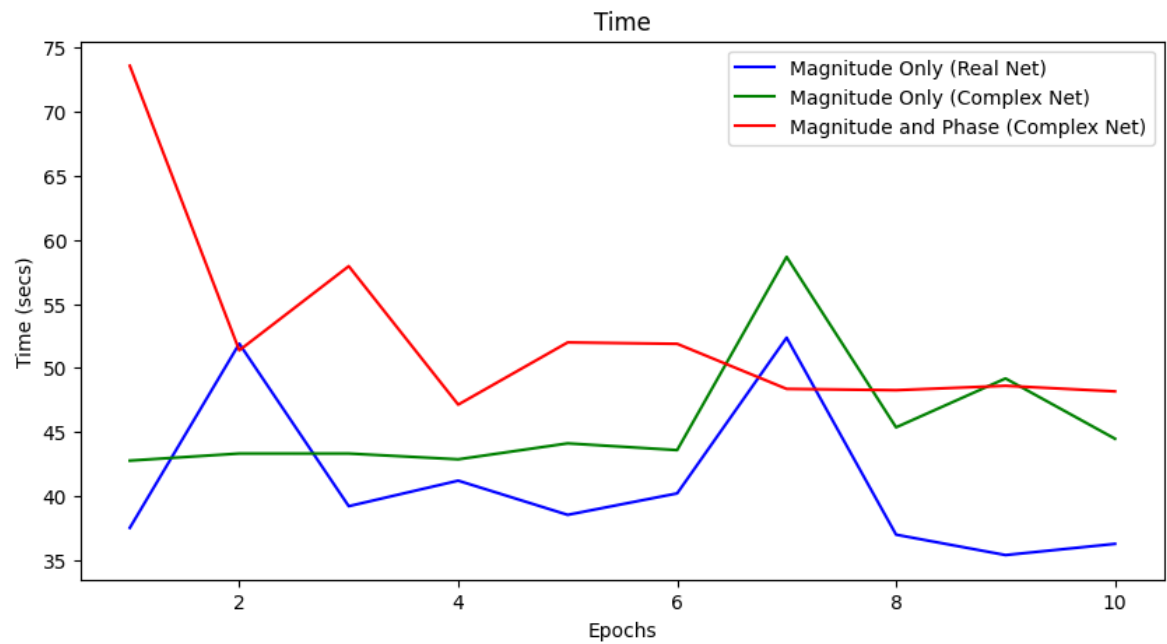
fig, axes = plt.subplots(1, 1, figsize=(10, 5))
for i, scenario in enumerate(scenarios):
    axes.plot(epochs, data[scenario]["epoch_times"], label=scenario, color=colors[i])
axes.set_title("Time")
axes.set_xlabel("Epochs")
axes.set_ylabel("Time (secs)")
axes.legend()

```



Out[35]: <matplotlib.legend.Legend at 0x2a9d28a90>



New way to Extract complex valued mfccs

```

In [52]: class MusicFeatureExtractorComplex2:
    def __init__(self, FFT_size=2048, HOP_SIZE=512, mel_filter_num=13, dct_filter_num=13, epsilon=1e-10):
        self.FFT_size = FFT_size
        self.HOP_SIZE = HOP_SIZE
        self.mel_filter_num = mel_filter_num
        self.dct_filter_num = dct_filter_num
        self.epsilon = 1e-10 # Added to log to avoid log10(0)

    def normalize_audio(self, audio):
        audio = audio / np.max(np.abs(audio))
        return audio

    def frame_audio(self, audio):
        frame_num = int((len(audio) - self.FFT_size) / self.HOP_SIZE) + 1
        frames = np.zeros((frame_num, self.FFT_size))
        for n in range(frame_num):
            frames[n] = audio[n * self.HOP_SIZE: n * self.HOP_SIZE + self.FFT_size]
        return frames

    def freq_to_mel(self, freq):
        return 2595.0 * np.log10(1.0 + freq / 700.0)

    def mel_to_freq(self, mels):
        return 700.0 * (10.0 ** (mels / 2595.0) - 1.0)

    def get_filter_points(self, fmin, fmax, sample_rate):
        fmin_mel = self.freq_to_mel(fmin)
        fmax_mel = self.freq_to_mel(fmax)
        mels = np.linspace(fmin_mel, fmax_mel, num=self.mel_filter_num + 1)
        freqs = self.mel_to_freq(mels)
        return np.floor((self.FFT_size + 1) / sample_rate * freqs).astype(int)

    def get_filters(self, filter_points):
        filters = np.zeros((len(filter_points) - 2, int(self.FFT_size / 2)))
        for n in range(len(filter_points) - 2):
            filters[n, filter_points[n]: filter_points[n + 1]] = np.linspace(0, 1, filter_points[n + 1] - filter_points[n])
            filters[n, filter_points[n + 1]: filter_points[n + 2]] = np.linspace(1, 0, filter_points[n + 2] - filter_points[n + 1])
        return filters

    def dct(self):
        basis = np.empty((self.dct_filter_num, self.mel_filter_num))
        basis[0, :] = 1.0 / np.sqrt(self.mel_filter_num)
        samples = np.arange(1, 2 * self.mel_filter_num, 2) * np.pi / (2.0 * self.FFT_size)
        for i in range(1, self.dct_filter_num):
            basis[i, :] = np.cos(i * samples) * np.sqrt(2.0 / self.mel_filter_num)
        return basis

    def get_mfcc_features(self, audio, sample_rate):
        audio = self.normalize_audio(audio)
        audio_framed = self.frame_audio(audio)
        window = get_window("hann", self.FFT_size, fftbins=True)
        audio_win = audio_framed * window
        audio_winT = np.transpose(audio_win)
        audio_fft = np.empty((int(1 + self.FFT_size // 2), audio_winT.shape[1]))
        for n in range(audio_fft.shape[1]):
            audio_fft[:, n] = fft.fft(audio_winT[:, n], axis=0)[:audio_fft.shape[0]]
        audio_fft = np.transpose(audio_fft)
        mag_fft = np.square(np.abs(audio_fft))
        phase_fft = np.angle(audio_fft)

```

```

freq_min = 0
freq_high = sample_rate / 2
filter_points, mel_freqs = self.get_filter_points(freq_min, freq_high)
filters = self.get_filters(filter_points)
audio_filtered = np.dot(filters, np.transpose(mag_fft))
phase_filtered = np.dot(filters, np.transpose(phase_fft))
audio_filtered = np.maximum(audio_filtered, self.epsilon) # Replace negative values with epsilon
audio_log = 10.0 * np.log10(audio_filtered)
dct_filters = self.dct()
cepstral_coefficients = np.dot(dct_filters, audio_log)
phase_coefficients = np.dot(dct_filters, phase_filtered)
return np.array([cepstral_coefficients*np.exp(1j*phase_coefficients)

```

```
class GenreDatasetPhaseMFCC2(GenreDatasetMFCC):
```

```

def __init__(self, train_path, n_fft=2048, hop_length=512, num_segments=10,
             super().__init__(train_path, n_fft, hop_length, num_segments, mel_filter_num=mel_filter_num)
self.mfcc_extractor = MusicFeatureExtractorComplex2(
    FFT_size=n_fft, HOP_SIZE=hop_length, mel_filter_num = mel_filter_num)

def __getitem__(self, idx):
    cur_file = self.files[idx]
    d = cur_file[0]
    file_path = cur_file[1]
    target = genre_mappings[str(file_path).split("/")[-1]]
    signal, sample_rate = librosa.load(file_path, sr=SAMPLE_RATE)
    start = self.samples_per_segment * d
    finish = start + self.samples_per_segment
    cur_signal = signal[start:finish]
    if self.training: cur_signal = self.apply_augmentations(cur_signal)
    cur_mfcc = self.mfcc_extractor.get_mfcc_features(cur_signal, sample_rate)
    cur_mfcc = self.adjust_shape(cur_mfcc)
    return torch.tensor(cur_mfcc, dtype=torch.complex64), target

```

```

In [53]: train_dataset = GenreDatasetPhaseMFCC2("Data/binary_data/train/", n_fft=2048, hop_length=512, num_segments=10)
test_dataset = GenreDatasetPhaseMFCC2("Data/binary_data/test/", n_fft=2048, hop_length=512, num_segments=10)
train_loader = torch.utils.data.DataLoader(dataset=train_dataset, shuffle=True)
test_loader = torch.utils.data.DataLoader(dataset=test_dataset, shuffle=False)

```

```
class ComplexNet(nn.Module):
```

```

def __init__(self):
    super(ComplexNet, self).__init__()
    self.conv1 = ComplexConv2d(1, 10, 2, 1)
    self.bn = ComplexBatchNorm2d(10)
    self.conv2 = ComplexConv2d(10, 20, 2, 1)
    self.fc1 = ComplexLinear(30*2*20, 500)
    self.fc2 = ComplexLinear(500, 2)

def forward(self, x):
    x = self.conv1(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = self.bn(x)
    x = self.conv2(x)
    x = complex_relu(x)
    x = complex_max_pool2d(x, 2, 2)
    x = x.view(-1, 30*2*20)
    x = self.fc1(x)
    x = complex_relu(x)

```

```
x = self.fc2(x)
x = x.abs()
x = F.log_softmax(x, dim=1)
return x

device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
model = ComplexNet().to(device)
optimizer = torch.optim.SGD(model.parameters(), lr=0.01, momentum=0.9)

metrics_dict_e4 = {
    'epoch_times': [],
    'train_losses': [],
    'train_accuracies': [],
    'test_losses': [],
    'test_accuracies': []
}

for epoch in range(NUM_EPOCHS):
    train(model,
          device,
          train_loader,
          test_loader,
          optimizer,
          epoch,
          metrics_dict_e4)

print("-"*100)
print("-"*100)
print("FINAL RESULTS:")
print("-"*100)
for key, value in metrics_dict_e3.items():
    print(f'{key}: {value}')
```

Train Epoch: 0 [0/ 1680 (0%)] Loss: 0.756448 Accuracy: 43.75%
Train Epoch: 0 [320/ 1680 (19%)] Loss: 0.943351 Accuracy: 51.70%
Train Epoch: 0 [640/ 1680 (38%)] Loss: 0.641432 Accuracy: 51.64%
Train Epoch: 0 [960/ 1680 (58%)] Loss: 0.497405 Accuracy: 56.96%
Train Epoch: 0 [1280/ 1680 (77%)] Loss: 0.510098 Accuracy: 61.36%
Train Epoch: 0 [1600/ 1680 (96%)] Loss: 0.745694 Accuracy: 62.99%
Epoch 0 – Time: 47.31s – Train Loss: 0.669079 – Train Accuracy: 62.74%
Test Loss: 0.466456 – Test Accuracy: 78.75%

Train Epoch: 1 [0/ 1680 (0%)] Loss: 0.427284 Accuracy: 81.25%
Train Epoch: 1 [320/ 1680 (19%)] Loss: 0.414837 Accuracy: 79.26%
Train Epoch: 1 [640/ 1680 (38%)] Loss: 0.425875 Accuracy: 77.83%
Train Epoch: 1 [960/ 1680 (58%)] Loss: 0.465073 Accuracy: 79.23%
Train Epoch: 1 [1280/ 1680 (77%)] Loss: 0.417217 Accuracy: 79.04%
Train Epoch: 1 [1600/ 1680 (96%)] Loss: 0.525346 Accuracy: 78.37%
Epoch 1 – Time: 49.40s – Train Loss: 0.464559 – Train Accuracy: 77.56%
Test Loss: 0.380537 – Test Accuracy: 81.56%

Train Epoch: 2 [0/ 1680 (0%)] Loss: 0.328772 Accuracy: 90.62%
Train Epoch: 2 [320/ 1680 (19%)] Loss: 0.275012 Accuracy: 79.55%
Train Epoch: 2 [640/ 1680 (38%)] Loss: 0.526225 Accuracy: 74.55%
Train Epoch: 2 [960/ 1680 (58%)] Loss: 0.534228 Accuracy: 78.12%
Train Epoch: 2 [1280/ 1680 (77%)] Loss: 0.513447 Accuracy: 79.04%
Train Epoch: 2 [1600/ 1680 (96%)] Loss: 0.389921 Accuracy: 79.66%
Epoch 2 – Time: 49.73s – Train Loss: 0.462108 – Train Accuracy: 78.93%
Test Loss: 0.376359 – Test Accuracy: 83.12%

Train Epoch: 3 [0/ 1680 (0%)] Loss: 0.300185 Accuracy: 87.50%
Train Epoch: 3 [320/ 1680 (19%)] Loss: 0.346790 Accuracy: 80.97%
Train Epoch: 3 [640/ 1680 (38%)] Loss: 0.440910 Accuracy: 81.10%
Train Epoch: 3 [960/ 1680 (58%)] Loss: 0.374527 Accuracy: 80.65%
Train Epoch: 3 [1280/ 1680 (77%)] Loss: 0.237164 Accuracy: 81.33%
Train Epoch: 3 [1600/ 1680 (96%)] Loss: 0.439932 Accuracy: 81.31%
Epoch 3 – Time: 48.36s – Train Loss: 0.390184 – Train Accuracy: 80.48%
Test Loss: 0.349970 – Test Accuracy: 86.25%

Train Epoch: 4 [0/ 1680 (0%)] Loss: 0.272466 Accuracy: 78.12%
Train Epoch: 4 [320/ 1680 (19%)] Loss: 0.492965 Accuracy: 85.80%
Train Epoch: 4 [640/ 1680 (38%)] Loss: 0.516339 Accuracy: 84.67%
Train Epoch: 4 [960/ 1680 (58%)] Loss: 0.494687 Accuracy: 84.98%
Train Epoch: 4 [1280/ 1680 (77%)] Loss: 0.296013 Accuracy: 86.05%
Train Epoch: 4 [1600/ 1680 (96%)] Loss: 0.349828 Accuracy: 86.27%
Epoch 4 – Time: 47.35s – Train Loss: 0.330548 – Train Accuracy: 85.54%
Test Loss: 0.290894 – Test Accuracy: 85.94%

Train Epoch: 5 [0/ 1680 (0%)] Loss: 0.229566 Accuracy: 87.50%
Train Epoch: 5 [320/ 1680 (19%)] Loss: 0.342616 Accuracy: 86.93%
Train Epoch: 5 [640/ 1680 (38%)] Loss: 0.348982 Accuracy: 87.95%
Train Epoch: 5 [960/ 1680 (58%)] Loss: 0.236092 Accuracy: 88.31%
Train Epoch: 5 [1280/ 1680 (77%)] Loss: 0.178407 Accuracy: 86.81%
Train Epoch: 5 [1600/ 1680 (96%)] Loss: 0.365560 Accuracy: 86.15%
Epoch 5 – Time: 51.33s – Train Loss: 0.331476 – Train Accuracy: 85.42%
Test Loss: 0.256633 – Test Accuracy: 90.62%

Train Epoch: 6 [0/ 1680 (0%)] Loss: 0.229851 Accuracy: 93.75%
Train Epoch: 6 [320/ 1680 (19%)] Loss: 0.202679 Accuracy: 88.92%
Train Epoch: 6 [640/ 1680 (38%)] Loss: 0.370229 Accuracy: 86.61%
Train Epoch: 6 [960/ 1680 (58%)] Loss: 0.450266 Accuracy: 86.49%
Train Epoch: 6 [1280/ 1680 (77%)] Loss: 0.280140 Accuracy: 85.52%
Train Epoch: 6 [1600/ 1680 (96%)] Loss: 0.274103 Accuracy: 85.17%

Epoch 6 – Time: 49.00s – Train Loss: 0.342386 – Train Accuracy: 84.52%
 Test Loss: 0.264726 – Test Accuracy: 90.31%

Train Epoch: 7 [0/ 1680 (0%)] Loss: 0.195107 Accuracy: 96.88%
 Train Epoch: 7 [320/ 1680 (19%)] Loss: 0.341948 Accuracy: 85.80%
 Train Epoch: 7 [640/ 1680 (38%)] Loss: 0.352814 Accuracy: 86.01%
 Train Epoch: 7 [960/ 1680 (58%)] Loss: 0.173795 Accuracy: 86.90%
 Train Epoch: 7 [1280/ 1680 (77%)] Loss: 0.332959 Accuracy: 86.51%
 Train Epoch: 7 [1600/ 1680 (96%)] Loss: 0.185387 Accuracy: 86.83%
 Epoch 7 – Time: 49.10s – Train Loss: 0.324903 – Train Accuracy: 85.95%
 Test Loss: 0.259553 – Test Accuracy: 87.81%

Train Epoch: 8 [0/ 1680 (0%)] Loss: 0.358597 Accuracy: 75.00%
 Train Epoch: 8 [320/ 1680 (19%)] Loss: 0.416254 Accuracy: 87.50%
 Train Epoch: 8 [640/ 1680 (38%)] Loss: 0.467524 Accuracy: 87.35%
 Train Epoch: 8 [960/ 1680 (58%)] Loss: 0.174863 Accuracy: 86.59%
 Train Epoch: 8 [1280/ 1680 (77%)] Loss: 0.237571 Accuracy: 87.27%
 Train Epoch: 8 [1600/ 1680 (96%)] Loss: 0.638832 Accuracy: 86.83%
 Epoch 8 – Time: 48.76s – Train Loss: 0.315972 – Train Accuracy: 85.83%
 Test Loss: 0.280964 – Test Accuracy: 85.62%

Train Epoch: 9 [0/ 1680 (0%)] Loss: 0.251841 Accuracy: 90.62%
 Train Epoch: 9 [320/ 1680 (19%)] Loss: 0.296167 Accuracy: 87.50%
 Train Epoch: 9 [640/ 1680 (38%)] Loss: 0.168065 Accuracy: 87.65%
 Train Epoch: 9 [960/ 1680 (58%)] Loss: 0.304702 Accuracy: 87.10%
 Train Epoch: 9 [1280/ 1680 (77%)] Loss: 0.210842 Accuracy: 87.20%
 Train Epoch: 9 [1600/ 1680 (96%)] Loss: 0.331602 Accuracy: 87.19%
 Epoch 9 – Time: 49.46s – Train Loss: 0.307399 – Train Accuracy: 86.43%
 Test Loss: 0.240680 – Test Accuracy: 90.94%

FINAL RESULTS:

```
epoch_times: [73.57843279838562, 51.3704469203949, 57.93836307525635, 47.1
35130882263184, 51.98855710029602, 51.87985110282898, 48.36233329772949, 4
8.26110076904297, 48.60330104827881, 48.16680717468262]
train_losses: [0.603170841645736, 0.4243808463215828, 0.30570677667856216,
0.285976423523747, 0.2851490078923794, 0.30193280887145263, 0.282665444108
1561, 0.2698039590882567, 0.25662878531819355, 0.25511147984518456]
train_accuracies: [71.30952380952381, 81.60714285714286, 86.1309523809523
8, 88.33333333333333, 87.5, 87.32142857142857, 88.33333333333333, 88.45238
095238095, 88.98809523809524, 88.21428571428571]
test_losses: [0.3362645523622632, 0.6025413427501917, 0.2809987593907863,
0.2766452480107546, 0.24619730543345214, 0.27177466712892057, 0.2100593810
9010458, 0.20127928256988525, 0.29388898848555983, 0.2773632241412997]
test_accuracies: [87.5, 66.5625, 90.3125, 91.5625, 90.9375, 89.0625, 90.93
75, 92.8125, 90.0, 91.5625]
```

```
In [55]: # Data for the four scenarios
data = {
    "Magnitude Only (Real Net)": metrics_dict_e1,
    "Magnitude Only (Complex Net)": metrics_dict_e2,
    "Magnitude and Phase (Complex Net)": metrics_dict_e3,
    "Magnitude and Phase (Complex Net), New Workflow": metrics_dict_e4
}
```

```

# Data for plotting
epochs = range(1, 11)
colors = ['b', 'g', 'r', 'm', 'y']
scenarios = list(data.keys())

fig, axes = plt.subplots(2, 1, figsize=(10, 10))

for i, scenario in enumerate(scenarios):
    axes[0].plot(epochs, data[scenario]["train_accuracies"], label=scenario)

axes[0].set_title("Train Accuracy")
axes[0].set_xlabel("Epochs")
axes[0].set_ylabel("Train Accuracy")
axes[0].legend()

for i, scenario in enumerate(scenarios):
    axes[1].plot(epochs, data[scenario]["test_accuracies"], label=scenario)

axes[1].set_title("Test Accuracy")
axes[1].set_xlabel("Epochs")
axes[1].set_ylabel("Test Accuracy")
axes[1].legend()

plt.tight_layout()
plt.show()

fig, axes = plt.subplots(2, 1, figsize=(10, 10))

for i, scenario in enumerate(scenarios):
    axes[0].plot(epochs, data[scenario]["train_losses"], label=scenario, color=colors[i])

axes[0].set_title("Train Loss")
axes[0].set_xlabel("Epochs")
axes[0].set_ylabel("Train Loss")
axes[0].legend()


for i, scenario in enumerate(scenarios):
    axes[1].plot(epochs, data[scenario]["test_losses"], label=scenario, color=colors[i])


axes[1].set_title("Test Loss")
axes[1].set_xlabel("Epochs")
axes[1].set_ylabel("Test Loss")
axes[1].legend()

plt.tight_layout()
plt.show()

fig, axes = plt.subplots(1, 1, figsize=(10, 5))
for i, scenario in enumerate(scenarios):
    axes.plot(epochs, data[scenario]["epoch_times"], label=scenario, color=colors[i])
axes.set_title("Time")
axes.set_xlabel("Epochs")
axes.set_ylabel("Time (secs)")
axes.legend()

```



Out[55]: <matplotlib.legend.Legend at 0x2b777f7d0>

